

▼ Problem 5 (Chapter 4, Exercise 10)

```
import pandas as pd
weekly = pd.read_csv("Weekly.csv")
```

```
weekly.info()
```

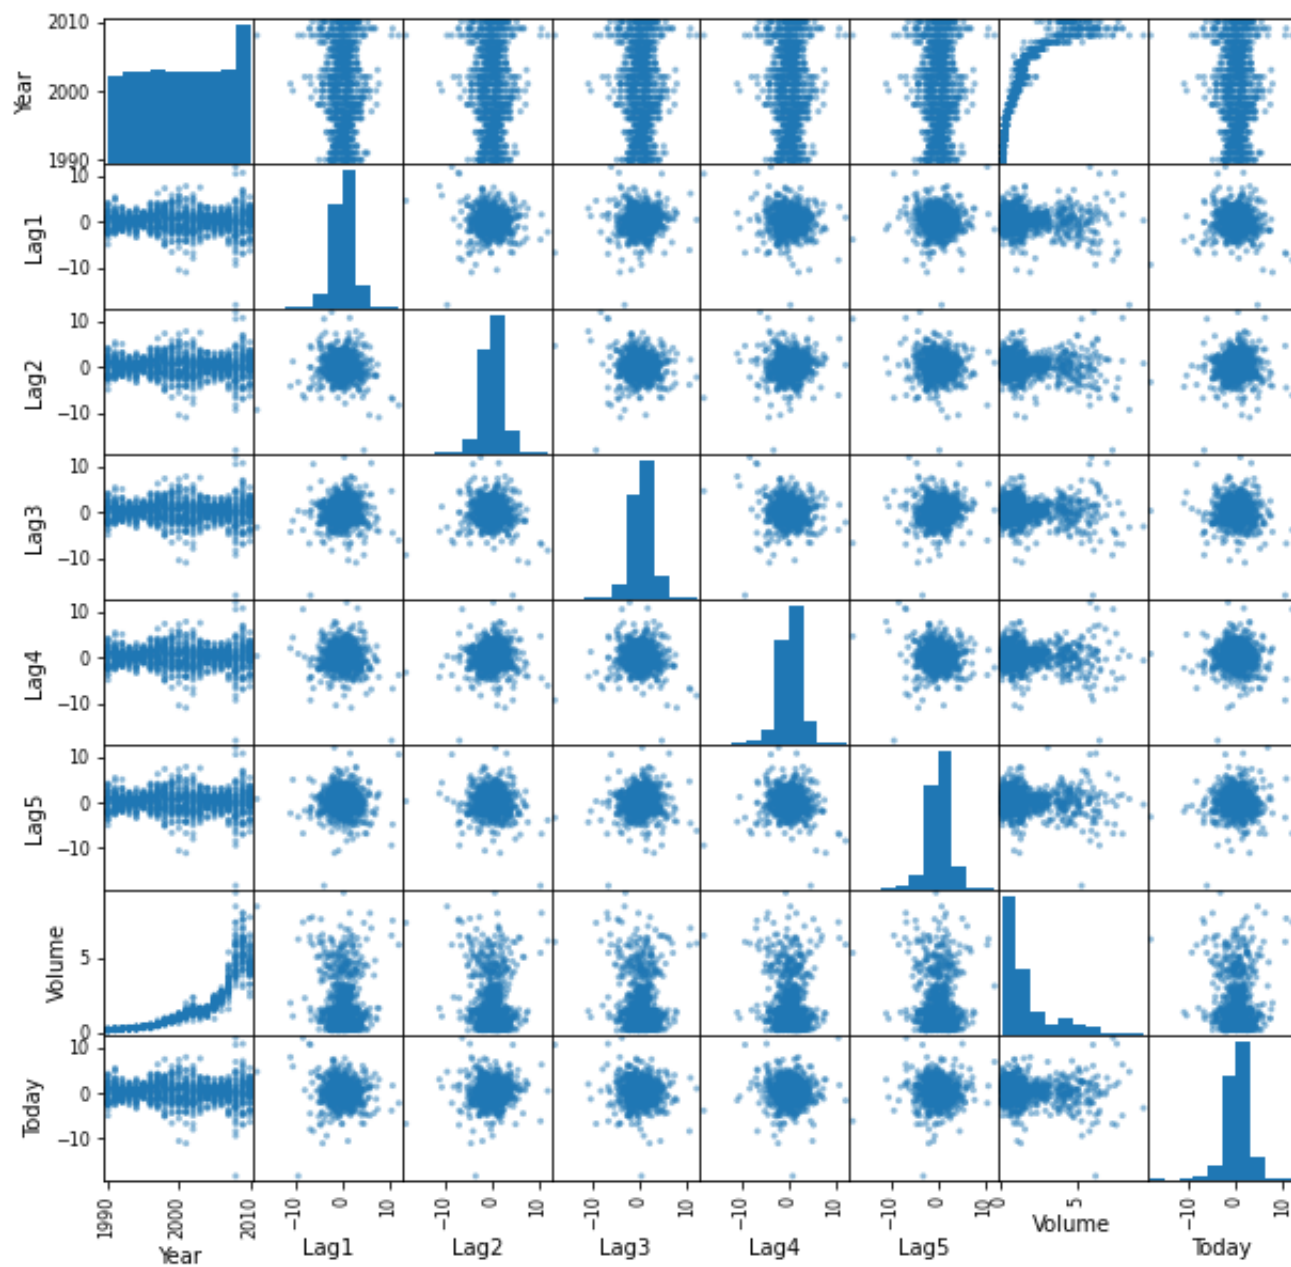
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1089 entries, 0 to 1088
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Year        1089 non-null   int64
1   Lag1        1089 non-null   float64
2   Lag2        1089 non-null   float64
3   Lag3        1089 non-null   float64
4   Lag4        1089 non-null   float64
5   Lag5        1089 non-null   float64
6   Volume      1089 non-null   float64
7   Today       1089 non-null   float64
8   Direction   1089 non-null   object
dtypes: float64(7), int64(1), object(1)
memory usage: 76.7+ KB
```

▼ Problem 5(a)

```
weekly.describe(include='all')
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	
count	1089.000000	1089.000000	1089.000000	1089.000000	1089.000000	1089.000000	1089.000000
unique	NaN	NaN	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	2000.048669	0.150585	0.151079	0.147205	0.145818	0.139893	0.139893
std	6.033182	2.357013	2.357254	2.360502	2.360279	2.361285	2.361285
min	1990.000000	-18.195000	-18.195000	-18.195000	-18.195000	-18.195000	-18.195000
25%	1995.000000	-1.154000	-1.154000	-1.158000	-1.158000	-1.166000	-1.166000
50%	2000.000000	0.241000	0.241000	0.241000	0.238000	0.234000	0.234000
75%	2005.000000	1.405000	1.409000	1.409000	1.409000	1.405000	1.405000
max	2010.000000	12.026000	12.026000	12.026000	12.026000	12.026000	12.026000

```
pd.plotting.scatter_matrix(weekly, figsize=(10,10));
```



```
weekly.corr()
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today
Year	1.000000	-0.032289	-0.033390	-0.030006	-0.031128	-0.030519	0.841942	-0.0324
Lag1	-0.032289	1.000000	-0.074853	0.058636	-0.071274	-0.008183	-0.064951	-0.0750
Lag2	-0.033390	-0.074853	1.000000	-0.075721	0.058382	-0.072499	-0.085513	0.0597
Lag3	-0.030006	0.058636	-0.075721	1.000000	-0.075396	0.060657	-0.069288	-0.0712
Lag4	-0.031128	-0.071274	0.058382	-0.075396	1.000000	-0.075675	-0.061075	-0.0078
Lag5	-0.030519	-0.008183	-0.072499	0.060657	-0.075675	1.000000	-0.058517	0.0110
Volume	0.841942	-0.064951	-0.085513	-0.069288	-0.061075	-0.058517	1.000000	-0.0330
Today	-0.032460	-0.075032	0.059167	-0.071244	-0.007826	0.011013	-0.033078	1.0000

▼ Problem 5(b)

```
import statsmodels.api as sm
```

```
X = weekly[['Lag1', 'Lag2', 'Lag3', 'Lag4', 'Lag5', 'Volume']]
```

```
X = sm.add_constant(X)
```

```
y = weekly['Direction'].astype('category').cat.codes
```

```
log_reg = sm.Logit(y, X).fit()
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning:
```

```
import pandas.util.testing as tm
```

```
Optimization terminated successfully.
```

```
Current function value: 0.682441
```

```
Iterations 4
```

```
print(log_reg.summary())
```

Logit Regression Results						
Dep. Variable:	y	No. Observations:	1089			
Model:	Logit	Df Residuals:	1087			
Method:	MLE	Df Model:	2			
Date:	Fri, 16 Jul 2021	Pseudo R-squ.:	0.006586			
Time:	19:36:57	Log-Likelihood:	-743.18			
converged:	True	LL-Null:	-748.16			
Covariance Type:	nonrobust	LLR p-value:	0.1313			
	coef	std err	z	P> z	[0.025	0.975]
const	0.2669	0.086	3.106	0.002	0.098	0.435
Lag1	-0.0413	0.026	-1.563	0.118	-0.093	0.010
Lag2	0.0584	0.027	2.175	0.030	0.006	0.110
Lag3	-0.0161	0.027	-0.602	0.547	-0.068	0.036
Lag4	-0.0278	0.026	-1.050	0.294	-0.080	0.024
Lag5	-0.0145	0.026	-0.549	0.583	-0.066	0.037
Volume	-0.0227	0.037	-0.616	0.538	-0.095	0.050

▼ Problem 5(c)

```
print(log_reg.pred_table())
```

```
[[ 54. 430.]
 [ 48. 557.]]
```

```
acc = (54 + 557) / (54 + 430 + 48 + 557)
print(acc)
```

```
0.5610651974288338
```

```
print(54/(54+430))
print(557/(557+48))
```

```
0.1115702479338843
0.9206611570247933
```

▼ Problem 5(d)

```

weekly1990to2008 = weekly.copy()
weekly1990to2008 = weekly1990to2008[weekly1990to2008['Year'] >= 1990]
weekly1990to2008 = weekly1990to2008[weekly1990to2008['Year'] <= 2008]
X_train = weekly1990to2008['Lag2']
X_train = sm.add_constant(X_train)

```

```

y_train = weekly1990to2008['Direction'].astype('category').cat.codes

```

```

log_reg = sm.Logit(y_train, X_train).fit()

```

```

    Optimization terminated successfully.
        Current function value: 0.685555
        Iterations 4

```

```

weekly2009to2010 = weekly.copy()
weekly2009to2010 = weekly2009to2010[weekly2009to2010['Year'] >= 2009]
weekly2009to2010 = weekly2009to2010[weekly2009to2010['Year'] <= 2010]
X_test = weekly2009to2010['Lag2']
X_test = sm.add_constant(X_test)

```

```

y_test = weekly2009to2010['Direction'].astype('category').cat.codes
y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

```

```

import numpy as np
cm = np.zeros((2,2))
for i in range(2):
    for j in range(2):
        cm[i, j] = sum(yt==i and yp==j for yt, yp in zip(y_test, y_pred))

```

```

cm

```

```

    array([[ 9., 34.],
           [ 5., 56.]])

```

```

acc = (9 + 56) / (9 + 56 + 5 + 34)
print(acc)

```

```

    0.625

```

```

print(9/(9+34))
print(56/(5+56))

```

```

    0.20930232558139536
    0.9180327868852459

```

▼ Problem 5(e)

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA()
lda.fit(X_train, y_train)

LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None,
                           solver='svd', store_covariance=False, tol=0.0001)

from sklearn.metrics import confusion_matrix

y_pred = lda.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)

[[ 9 34]
 [ 5 56]]

acc = (9 + 56) / (9 + 56 + 5 + 34)
print(acc)

0.625

print(9/(9+34))
print(56/(5+56))

0.20930232558139536
0.9180327868852459
```

▼ Problem 5(f)

[illegible]

```

y_pred = qda.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)

[[43  0]
 [61  0]]
/usr/local/lib/python3.7/dist-packages/sklearn/discriminant_analysis.py:715: F
  X2 = np.dot(Xm, R * (S ** (-0.5)))
/usr/local/lib/python3.7/dist-packages/sklearn/discriminant_analysis.py:715: F
  X2 = np.dot(Xm, R * (S ** (-0.5)))
/usr/local/lib/python3.7/dist-packages/sklearn/discriminant_analysis.py:718: F
  u = np.asarray([np.sum(np.log(s)) for s in self.scalings_])

acc = (43) / (43+61)
print(acc)

0.41346153846153844

print(43/(43+0))
print(0/(61+0))

1.0
0.0

```

▼ Problem 5(g)

```

from sklearn.neighbors import KNeighborsClassifier as KNN

knn = KNN(1)
knn.fit(X_train, y_train)

KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=1, p=2,
                    weights='uniform')

y_pred = knn.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)

[[21 22]
 [31 30]]

```



```
acc = (21+30) / (21+30+22+31)
print(acc)
```

```
0.49038461538461536
```

```
print(21/(21+22))
print(30/(31+30))
```


```
0.4883720930232558
```

```
0.4918032786885246
```

▼ Problem 6 (Chapter 5, Exercise 2)

▼ Problem 6(g)

```
import matplotlib.pyplot as plt
plt.plot(np.arange(1, 100000), [1-(1-1/n)**n for n in np.arange(1, 100000)])
plt.xlabel("n")
plt.ylabel("P(x_j is in X^b)");
```



▼ Problem 6(h)

```
np.random.seed(1)

ctr = 0
for i in range(10000):
    a = np.random.choice(range(100), size=100, replace=True)
    if 4 in a:
        ctr += 1

print(ctr)

6302
```

▼ Problem 7 (Chapter 5, Exercise 5)

```
default = pd.read_csv("Default.csv")

np.random.seed(1)
```

▼ Problem 7(a)

```
X = default[['income', 'balance']]
X = sm.add_constant(X)

y = default['default'].astype('category').cat.codes

log_reg = sm.Logit(y, X).fit()

Optimization terminated successfully.
    Current function value: 0.078948
    Iterations 10
```

▼ Problem 7(b)

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y)

log_reg = sm.Logit(y_train, X_train).fit()

y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

error = np.mean([yt != yp for yt, yp in zip(y_test, y_pred)])
print(error)

Optimization terminated successfully.
      Current function value: 0.078638
      Iterations 10
0.0252
```

▼ Problem 7(c)

```

# trial 1
X_train, X_test, y_train, y_test = train_test_split(X, y)

log_reg = sm.Logit(y_train, X_train).fit()

y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

error = np.mean([yt != yp for yt, yp in zip(y_test, y_pred)])
print(error)

# trial 2
X_train, X_test, y_train, y_test = train_test_split(X, y)

log_reg = sm.Logit(y_train, X_train).fit()

y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

error = np.mean([yt != yp for yt, yp in zip(y_test, y_pred)])
print(error)

# trial 3
X_train, X_test, y_train, y_test = train_test_split(X, y)

log_reg = sm.Logit(y_train, X_train).fit()

y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

error = np.mean([yt != yp for yt, yp in zip(y_test, y_pred)])
print(error)

    Optimization terminated successfully.
        Current function value: 0.079030
        Iterations 10
0.0236
    Optimization terminated successfully.
        Current function value: 0.076866
        Iterations 10
0.0276
    Optimization terminated successfully.
        Current function value: 0.079056
        Iterations 10
0.0272

```

▼ Problem 7(d)

```

default['student'] = default['student'].astype('category').cat.codes

X = default[['income', 'balance', 'student']]
X = sm.add_constant(X)

X_train, X_test, y_train, y_test = train_test_split(X, y)

log_reg = sm.Logit(y_train, X_train).fit()

y_pred = (log_reg.predict(X_test) >= 0.5).astype(int)

error = np.mean([yt != yp for yt, yp in zip(y_test, y_pred)])
print(error)

Optimization terminated successfully.
      Current function value: 0.076317
      Iterations 10
0.028

```

▼ Problem 8 (Chapter 5, Exercise 6)

```

default = pd.read_csv("Default.csv")

np.random.seed(1)

```

▼ Problem 8(a)

```

X = default[['income', 'balance']]
X = sm.add_constant(X)

y = default['default'].astype('category').cat.codes

log_reg = sm.Logit(y, X).fit()

Optimization terminated successfully.
      Current function value: 0.078948
      Iterations 10

```

```
print(log_reg.summary())
```

Logit Regression Results						
Dep. Variable:	y	No. Observations:	10000			
Model:	Logit	Df Residuals:	9997			
Method:	MLE	Df Model:	3			
Date:	Fri, 16 Jul 2021	Pseudo R-squ.:	0.4594			
Time:	19:36:59	Log-Likelihood:	-789.48			
converged:	True	LL-Null:	-1460.3			
Covariance Type:	nonrobust	LLR p-value:	4.541e-29			
	coef	std err	z	P> z	[0.025	0.975]
const	-11.5405	0.435	-26.544	0.000	-12.393	-10.688
income	2.081e-05	4.99e-06	4.174	0.000	1.1e-05	3.06e-05
balance	0.0056	0.000	24.835	0.000	0.005	0.006

Possibly complete quasi-separation: A fraction 0.14 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

▼ Problem 8(b)

```
def boot_fn(default, indices):

    X = default[['income', 'balance']]
    X = X.iloc[indices, :]
    X = sm.add_constant(X)

    y = default['default'].astype('category').cat.codes
    y = y.iloc[indices]

    log_reg = sm.Logit(y, X).fit()

    return log_reg.params[1:]
```

▼ Problem 8(c)

```
beta_is, beta_bs = [], []
B = 100

for i in range(B):
```

```
beta_is.append(beta_b)
beta_bs.append(beta_b)
```

```
Current function value: 0.070837
Iterations 10
Optimization terminated successfully.
Current function value: 0.075469
Iterations 10
Optimization terminated successfully.
Current function value: 0.074056
Iterations 10
Optimization terminated successfully.
Current function value: 0.080399
Iterations 10
Optimization terminated successfully.
Current function value: 0.079418
Iterations 10
Optimization terminated successfully.
Current function value: 0.075494
Iterations 10
Optimization terminated successfully.
Current function value: 0.076585
Iterations 10
Optimization terminated successfully.
Current function value: 0.078997
Iterations 10
Optimization terminated successfully.
Current function value: 0.077621
Iterations 10
Optimization terminated successfully.
Current function value: 0.077644
Iterations 10
Optimization terminated successfully.
Current function value: 0.078250
Iterations 10
Optimization terminated successfully.
Current function value: 0.076597
Iterations 10
Optimization terminated successfully.
Current function value: 0.082268
Iterations 10
Optimization terminated successfully.
Current function value: 0.079406
Iterations 10
Optimization terminated successfully.
Current function value: 0.080869
Iterations 10
Optimization terminated successfully.
Current function value: 0.075792
Iterations 10
Optimization terminated successfully.
Current function value: 0.079341
Iterations 10
```

```
Optimization terminated successfully.
```

```

optimization terminated successfully.
    Current function value: 0.076661
    Iterations 10
Optimization terminated successfully.
    Current function value: 0.086022
    Iterations 10
Optimization terminated successfully.
    Current function value: 0.083204
    Iterations 10

```

```

beta_i_mean = np.mean(beta_is)
beta_i_se    = np.sqrt(1/(B-1)*np.sum( [(beta_i - beta_i_mean)**2 for beta_i in be

beta_b_mean = np.mean(beta_bs)
beta_b_se    = np.sqrt(1/(B-1)*np.sum( [(beta_b - beta_b_mean)**2 for beta_b in be

print("beta_income (mean & SE)", beta_i_mean, beta_i_se)
print("beta_balance (mean & SE)", beta_b_mean, beta_b_se)

beta_income (mean & SE) 2.016973494321125e-05 5.4383662396369285e-06
beta_balance (mean & SE) 0.005677686072213888 0.00022069873134705593

```

▼ Problem 9 (Chapter 5, Exercise 8)

▼ Problem 9(a)

```
data = pd.read_csv("ch5_ex8.csv")
```

```
data.info()
```


```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    x      100 non-null     float64
 1    y      100 non-null     float64
dtypes: float64(2)
memory usage: 1.7 KB

```

▼ Problem 9(b)


```
plt.scatter(data.x, data.y)
plt.xlabel("x")
plt.ylabel("y");
```



▼ Problem 9(c)

```
np.random.seed(1)
```

```

from sklearn.preprocessing import PolynomialFeatures

X = np.array(data['x']).reshape(-1, 1)
y = np.array(data['y']).reshape(-1, 1)

for k in [1, 2, 3, 4]:

    loocv_errors = []

    for i in range(data.shape[0]):

        X_loo = np.delete(X, i).reshape(-1, 1)
        y_loo = np.delete(y, i).reshape(-1, 1)

        pf = PolynomialFeatures(degree=k)
        Xp_loo = pf.fit_transform(X_loo)

        lin_reg = sm.OLS(y_loo, Xp_loo).fit()
        y_pred = lin_reg.predict(pf.transform(X[i].reshape(1,1)))
        loocv_error = (y_pred - y[i])**2

    loocv_errors.append(loocv_error)

loocv_errors = np.array(loocv_errors)
print("degree %d L00CV error:" %k, "%+3.2e" %loocv_errors.mean())

degree 1 L00CV error: +5.89e+00
degree 2 L00CV error: +1.09e+00
degree 3 L00CV error: +1.10e+00
degree 4 L00CV error: +1.11e+00

```

▼ Problem 9(d)

```
np.random.seed(2)
```

```

for k in [1, 2, 3, 4]:

    loocv_errors = []

    for i in range(data.shape[0]):

        X_loo = np.delete(X, i).reshape(-1, 1)
        y_loo = np.delete(y, i).reshape(-1, 1)

        pf = PolynomialFeatures(degree=k)
        Xp_loo = pf.fit_transform(X_loo)

        lin_reg = sm.OLS(y_loo, Xp_loo).fit()
        y_pred = lin_reg.predict(pf.transform(X[i].reshape(1,1)))
        loocv_error = (y_pred - y[i])**2

    loocv_errors.append(loocv_error)

loocv_errors = np.array(loocv_errors)
print("degree %d L00CV error:" %k, "%+3.2e" %loocv_errors.mean())

degree 1 L00CV error: +5.89e+00
degree 2 L00CV error: +1.09e+00
degree 3 L00CV error: +1.10e+00
degree 4 L00CV error: +1.11e+00

```

▼ Problem 9(f)

```

X = np.array(data['x']).reshape(-1, 1)
y = np.array(data['y']).reshape(-1, 1)

for k in [1, 2, 3, 4]:

    pf = PolynomialFeatures(degree=k)
    Xp = pf.fit_transform(X)

    lin_reg = sm.OLS(y, Xp).fit()
    print(lin_reg.summary())

```

OLS Regression Results

Dep. Variable:	y	R-squared:	0.010
Model:	OLS	Adj. R-squared:	-0.000
Method:	Least Squares	F-statistic:	0.9610
Date:	Fri, 16 Jul 2021	Prob (F-statistic):	0.329
Time:	19:37:03	Log-Likelihood:	-226.84
No. Observations:	100	AIC:	457.7

```

NO. OBSERVATIONS:      100      AIC:      457.1
Df Residuals:          98      BIC:      462.9
Df Model:              1
Covariance Type:      nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const        -1.8185      0.236      -7.692      0.000      -2.288      -1.349
x1           0.2430      0.248       0.981      0.329      -0.249      0.735
=====
Omnibus:      17.572      Durbin-Watson:      2.198
Prob(Omnibus): 0.000      Jarque-Bera (JB):      20.497
Skew:         -1.051      Prob(JB):      3.55e-05
Kurtosis:     3.704      Cond. No.      1.06
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correct
      OLS Regression Results

```

```

=====
Dep. Variable:      y      R-squared:      0.813
Model:              OLS      Adj. R-squared:      0.809
Method:              Least Squares      F-statistic:      210.6
Date:                Fri, 16 Jul 2021      Prob (F-statistic):      5.10e-36
Time:                19:37:03      Log-Likelihood:      -143.59
No. Observations:    100      AIC:      293.1
Df Residuals:        97      BIC:      300.9
Df Model:            2
Covariance Type:     nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const        -0.0954      0.133      -0.715      0.476      -0.360      0.169
x1           0.8996      0.113       7.961      0.000      0.675      1.124
x2          -1.8666      0.092     -20.399      0.000      -2.048     -1.685
=====
Omnibus:      1.794      Durbin-Watson:      2.236
Prob(Omnibus): 0.408      Jarque-Bera (JB):      1.229
Skew:         -0.183      Prob(JB):      0.542
Kurtosis:     3.399      Cond. No.      2.47
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correct
      OLS Regression Results

```

```

=====
Dep. Variable:      y      R-squared:      0.813
Model:              OLS      Adj. R-squared:      0.807
Method:              Least Squares      F-statistic:      139.1
Date:                Fri, 16 Jul 2021      Prob (F-statistic):      8.04e-35
Time:                19:37:03      Log-Likelihood:      -143.59
No. Observations:    100      AIC:      295.6
Df Residuals:        98      BIC:      305.1
Df Model:            2
Covariance Type:     nonrobust
=====

```

▼ Problem 10 (Chapter 5, Exercise 9)

```
boston = pd.read_csv("Boston.csv")
```

▼ Problem 10(a)

```
mu_hat = boston['medv'].mean()  
print(mu_hat)
```

```
22.532806324110698
```

▼ Problem 10(b)

```
SE_mu_hat = boston['medv'].std()/np.sqrt(boston.shape[0])  
print(SE_mu_hat)
```

```
0.4088611474975351
```

▼ Problem 10(c)

▼ Problem 10(d)

▼ Problem 10(e)

▼ Problem 10(f)

▼ Problem 10(g)

▼ Problem 10(h)