


### Homework 3


Ross B. Alexander (rbalexan@stanford.edu)

#### Problem 1 (C6E4)

(a)



(b)



## Problem 2 (C8E8)

(a) - (c)

See attached code.

### Problem 3 (CSE10)

(a) - (g) See attached code.

## Problem 4 (CSE11)

(a) - (c) See attached code.

## Problem 5

- (a)  $x^{(1)} : i = 1, \dots, P$  are predictor values  
 $a_k^{(2s)} : k = 1, \dots, K$  are  $K$ -dim. output from a 2-layer,  $M$ -hidden unit NN  
with sigmoid activation  $\sigma(a) = 1/(1+e^{-a})$  such that

$$a_j^{(1s)} = w_{j0}^{(1s)} + \sum_{i=1}^P w_{ji}^{(1s)} x_i \quad j = 1, \dots, M$$

$$a_k^{(2s)} = w_{k0}^{(2s)} + \sum_{j=1}^M w_{kj}^{(2s)} \sigma(a_j^{(1s)})$$

Show that there exists an equivalent network that computes exactly the same output values, but with hidden unit activations given by

$$\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}, \text{ i.e.}$$

$$a_j^{(1t)} = w_{j0}^{(1t)} + \sum_{i=1}^P w_{ji}^{(1t)} x_i \quad j = 1, \dots, M$$

$$a_k^{(2t)} = w_{k0}^{(2t)} + \sum_{j=1}^M w_{kj}^{(2t)} \tanh(a_j^{(1t)})$$

- we want to find a relationship between  $\sigma(a)$  and  $\tanh(a)$

$$\begin{aligned} \tanh(a) &= \frac{e^a - e^{-a}}{e^a + e^{-a}} \\ &= \frac{e^a - e^{-a} + (e^{-a} - e^{-a})}{e^a + e^{-a}} \\ &= \frac{(e^a + e^{-a}) - 2e^{-a}}{e^a + e^{-a}} \\ &= \frac{e^a + e^{-a}}{e^a + e^{-a}} - \frac{2e^{-a}}{e^a + e^{-a}} \\ &= 1 - 2 \left( \frac{e^{-a}}{e^a + e^{-a}} \right)^{-1} \\ &= 1 - 2 \left( e^{2a} + e^0 \right)^{-1} \\ &= 1 - \frac{2}{1 + e^{2a}} \end{aligned}$$

- leveraging our definition of  $\sigma(a)$

$$\tanh(a) = 1 - \frac{2}{1 + e^{2a}}$$

$$\tanh(a) = 1 - 2\sigma(-2a)$$

- substituting our relationship into the NN

$$a_j^{(1t)} = \omega_{j0}^{(1t)} + \sum_{i=1}^P \omega_{ji}^{(1t)} x_i \quad j=1, \dots, M$$

$$a_k^{(2t)} = \omega_{k0}^{(2t)} + \sum_{j=1}^M \omega_{kj}^{(2t)} (1 - 2\sigma(-2a_j^{(1t)}))$$

- expanding the summation and moving the  $\sigma(-2a_j^{(1t)})$  to  $\sigma(a_j^{(1t)})$  by multiplying  $a_j^{(1t)}$  by  $-2$ , we have

$$a_j^{(1t)} = -2 \left[ \omega_{j0}^{(1t)} + \sum_{i=1}^P \omega_{ji}^{(1t)} x_i \right] \quad j=1, \dots, M$$

$$a_k^{(2t)} = \omega_{k0}^{(2t)} + \sum_{j=1}^M \omega_{kj}^{(2t)} - 2 \sum_{j=1}^M \omega_{kj}^{(2t)} \sigma(a_j^{(1t)})$$

- in total, we now have

$$a_j^{(1t)} = -2\omega_{j0}^{(1t)} + \sum_{i=1}^P (-2\omega_{ji}^{(1t)}) x_i \quad j=1, \dots, M$$

$$a_k^{(2t)} = \sum_{j=0}^M \omega_{kj}^{(2t)} + \sum_{j=1}^M (-2\omega_{kj}^{(2t)}) \sigma(a_j^{(1t)})$$

- therefore the direct correspondance is

$\omega_{j0}^{(1s)} = -2\omega_{j0}^{(1t)}$	$\omega_{ji}^{(1s)} = -2\omega_{ji}^{(1t)}$
$\omega_{k0}^{(2s)} = \sum_{j=0}^M \omega_{kj}^{(2t)}$	$\omega_{kj}^{(2s)} = -2\omega_{kj}^{(2t)}$

```
In [1]: import numpy as np
import pandas as pd
```

## Problem 2 (Chapter 8, Exercise 8)

```
In [2]: carseats = pd.read_csv("data/Carseats.csv")

print(carseats.info())
carseats.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Sales           400 non-null   float64
1   CompPrice       400 non-null   int64   
2   Income         400 non-null   int64   
3   Advertising     400 non-null   int64   
4   Population      400 non-null   int64   
5   Price          400 non-null   int64   
6   ShelfLoc       400 non-null   object  
7   Age            400 non-null   int64   
8   Education       400 non-null   int64   
9   Urban          400 non-null   object  
10  US             400 non-null   object  
dtypes: float64(1), int64(7), object(3)
memory usage: 34.5+ KB
None
```

```
Out[2]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban
0	9.50	138	73	11	276	120	Bad	42	17	Ye
1	11.22	111	48	16	260	83	Good	65	10	Ye
2	10.06	113	35	10	269	80	Medium	59	12	Ye
3	7.40	117	100	4	466	97	Medium	55	14	Ye
4	4.15	141	64	3	340	128	Bad	38	13	Ye

```
In [3]: from sklearn.preprocessing import OneHotEncoder

# fix categorical columns
for cat in ["Urban", "US"]:
    carseats[cat] = carseats[cat].astype('category').cat.codes

enc = OneHotEncoder(sparse=False)
shelve_loc = enc.fit_transform(carseats["ShelveLoc"].to_numpy().reshape(-1, 1))

carseats["ShelveLocBad"] = shelve_loc[:, 0]
carseats["ShelveLocMedium"] = shelve_loc[:, 1]
carseats["ShelveLocGood"] = shelve_loc[:, 2]

carseats = carseats.drop("ShelveLoc", axis=1)

print(carseats.info())
carseats.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Sales                 400 non-null   float64
1   CompPrice             400 non-null   int64
2   Income               400 non-null   int64
3   Advertising           400 non-null   int64
4   Population            400 non-null   int64
5   Price                400 non-null   int64
6   Age                  400 non-null   int64
7   Education             400 non-null   int64
8   Urban                400 non-null   int8
9   US                   400 non-null   int8
10  ShelveLocBad          400 non-null   float64
11  ShelveLocMedium       400 non-null   float64
12  ShelveLocGood         400 non-null   float64
dtypes: float64(4), int64(7), int8(2)
memory usage: 35.3 KB
None
```

```
Out[3]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	Age	Education	Urban	US	She
0	9.50	138	73	11	276	120	42	17	1	1	
1	11.22	111	48	16	260	83	65	10	1	1	
2	10.06	113	35	10	269	80	59	12	1	1	
3	7.40	117	100	4	466	97	55	14	1	1	
4	4.15	141	64	3	340	128	38	13	1	0	

## Problem 2(a)



```
In [4]: from sklearn.model_selection import train_test_split

X = carseats.drop("Sales", axis=1)
y = carseats["Sales"]

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```


## Problem 2(b)

```
In [5]: from sklearn.tree import DecisionTreeRegressor, plot_tree

reg_tree = DecisionTreeRegressor(random_state=0)
reg_tree.fit(X_train, y_train);
```

```
In [6]: from matplotlib import pyplot as plt

fig = plt.figure(figsize=(15,15))
ax = fig.gca()
plot_tree(reg_tree, ax=ax, feature_names=carseats.drop("Sales", axis=1).columns,
          class_names="Sales", impurity=False);
```



The tree is very highly branched and possibly overfit to the dataset, so we may want to do some pruning.

```
In [7]: from sklearn.metrics import mean_squared_error

mse = mean_squared_error(y_test, reg_tree.predict(X_test))
print("decision tree test mse: %.3f" % mse)
```

```
decision tree test mse: 5.160
```

## Problem 2(c)

```
In [8]: from sklearn.model_selection import cross_validate

min_decr_list = np.logspace(-3, 3, 13)

for min_decr in min_decr_list:

    reg_tree = DecisionTreeRegressor(random_state=0, min_impurity_decrease=min_decr)

    cv_results = cross_validate(reg_tree, X, y, cv=10, scoring='neg_mean_squared_error')

    avg_test_mse = np.mean(-1*cv_results['test_score'])

    print("min split decr: %5.3e" % min_decr, " | test mse:", avg_test_mse)
```

min split decr: 1.000e-03	test mse: 5.021373217986111
min split decr: 3.162e-03	test mse: 4.878396738488783
min split decr: 1.000e-02	test mse: 4.661581790056902
min split decr: 3.162e-02	test mse: 4.418206140487827
min split decr: 1.000e-01	test mse: 4.756998382990023
min split decr: 3.162e-01	test mse: 4.903125250448153
min split decr: 1.000e+00	test mse: 6.114254875860025
min split decr: 3.162e+00	test mse: 8.024303693441357
min split decr: 1.000e+01	test mse: 8.024303693441357
min split decr: 3.162e+01	test mse: 8.024303693441357
min split decr: 1.000e+02	test mse: 8.024303693441357
min split decr: 3.162e+02	test mse: 8.024303693441357
min split decr: 1.000e+03	test mse: 8.024303693441357

```
In [9]: max_nodes_list = np.logspace(np.log10(2), np.log10(1000), 20)

for max_nodes in max_nodes_list:

    reg_tree = DecisionTreeRegressor(random_state=0, max_leaf_nodes=round(max_nodes))

    cv_results = cross_validate(reg_tree, X, y, cv=10, scoring='neg_mean_squared_error')

    avg_test_mse = np.mean(-1*cv_results['test_score'])

    print("max leaf nodes: %4d" % round(max_nodes), " | test mse:", avg_test_mse)
```

max leaf nodes:	2	test mse:	6.021211549802172
max leaf nodes:	3	test mse:	5.463939411950188
max leaf nodes:	4	test mse:	5.350477382126139
max leaf nodes:	5	test mse:	4.7615529635906455
max leaf nodes:	7	test mse:	4.869271383110496
max leaf nodes:	10	test mse:	4.74810583970232
max leaf nodes:	14	test mse:	4.80136840944212
max leaf nodes:	20	test mse:	4.639243894469821
max leaf nodes:	27	test mse:	4.62718150075911
max leaf nodes:	38	test mse:	4.415070541600826
max leaf nodes:	53	test mse:	4.466614060143851
max leaf nodes:	73	test mse:	4.552412057504686
max leaf nodes:	101	test mse:	4.759964075058629
max leaf nodes:	141	test mse:	4.94207044564928
max leaf nodes:	195	test mse:	5.063093958541667
max leaf nodes:	270	test mse:	5.071329741319444
max leaf nodes:	375	test mse:	5.07338375
max leaf nodes:	520	test mse:	5.07338375
max leaf nodes:	721	test mse:	5.07338375
max leaf nodes:	1000	test mse:	5.07338375

```
In [10]: max_depth_list = np.logspace(np.log10(2), np.log10(100), 20)

for max_depth in max_depth_list:

    reg_tree = DecisionTreeRegressor(random_state=0, max_depth=round(max_depth))

    cv_results = cross_validate(reg_tree, X, y, cv=10, scoring='neg_mean_squared_error')

    avg_test_mse = np.mean(-1*cv_results['test_score'])

    print("max tree depth: %3d" % round(max_depth), " | test mse:", avg_test_mse)
```

max tree depth:	2	test mse:	5.175545607747162
max tree depth:	2	test mse:	5.175545607747162
max tree depth:	3	test mse:	4.7590242840611285
max tree depth:	4	test mse:	4.882876270347563
max tree depth:	5	test mse:	4.646330628189368
max tree depth:	6	test mse:	4.5029336227383485
max tree depth:	7	test mse:	4.625605916258284
max tree depth:	8	test mse:	4.601803437534455
max tree depth:	10	test mse:	5.0851134190414635
max tree depth:	13	test mse:	5.103489814687501
max tree depth:	16	test mse:	5.111214562500001
max tree depth:	19	test mse:	4.9669585
max tree depth:	24	test mse:	4.9669585
max tree depth:	29	test mse:	4.9669585
max tree depth:	36	test mse:	4.9669585
max tree depth:	44	test mse:	4.9669585
max tree depth:	54	test mse:	4.9669585
max tree depth:	66	test mse:	4.9669585
max tree depth:	81	test mse:	4.9669585
max tree depth:	100	test mse:	4.9669585

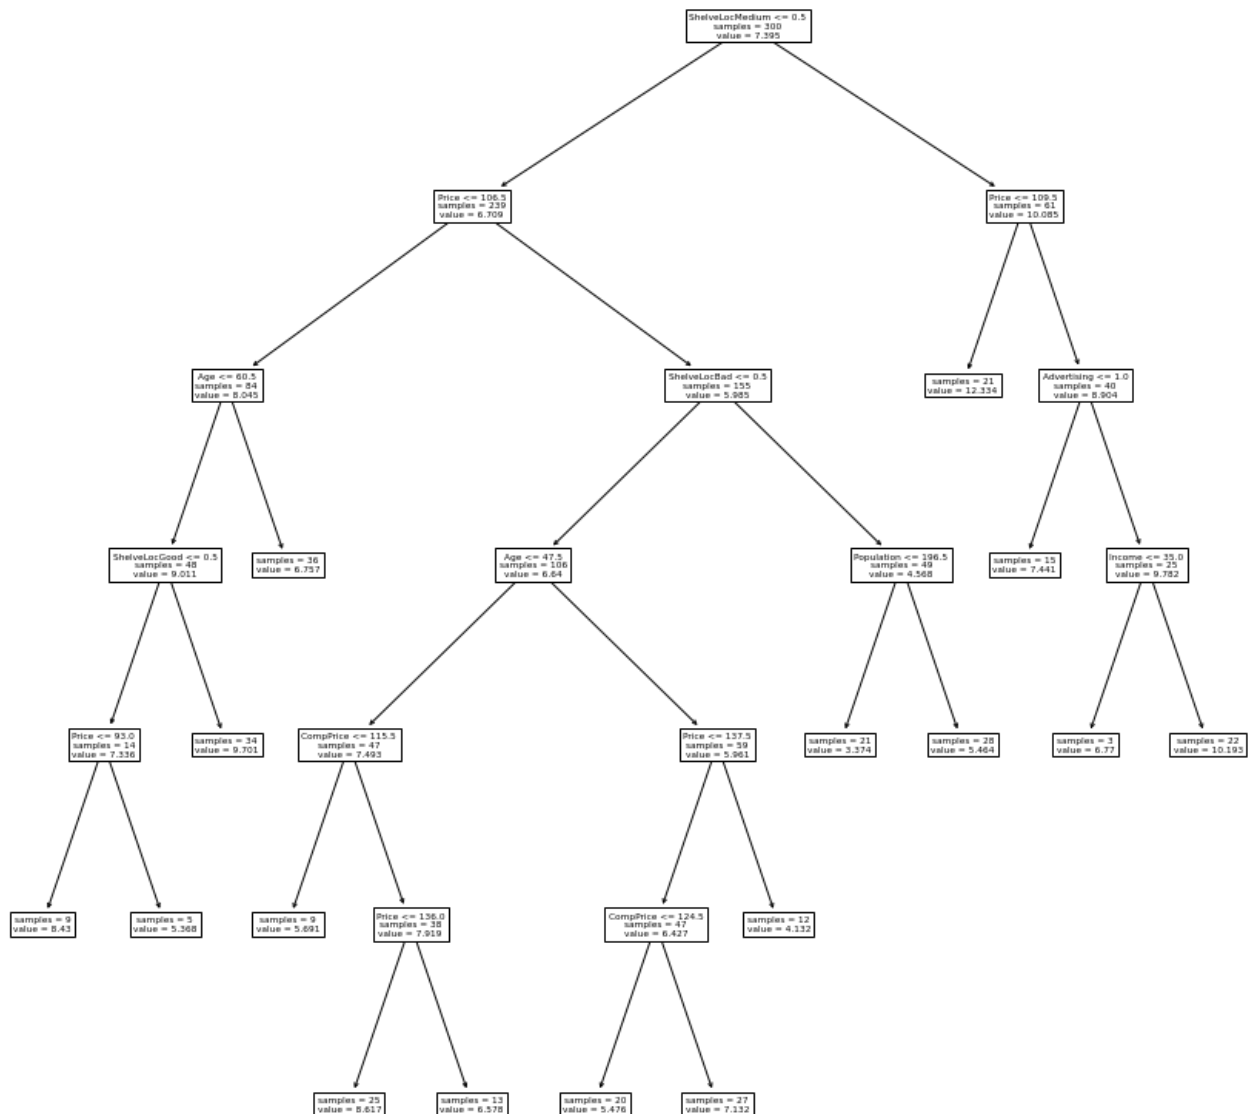
The optimal tree complexity is with min split decrease of  $\sim 3 \cdot 10^{-2}$ , max leaf nodes of  $\sim 50$ , and max tree depth of  $\sim 6$ .

```
In [11]: reg_tree = DecisionTreeRegressor(random_state=0, min_impurity_decrease=1E-1,
reg_tree.fit(X_train, y_train)

fig = plt.figure(figsize=(15,15))
ax = fig.gca()
plot_tree(reg_tree, ax=ax, feature_names=carseats.drop("Sales", axis=1).columns,
class_names="Sales", impurity=False);

mse = mean_squared_error(y_test, reg_tree.predict(X_test))
print("decision tree test mse: %.3f" % mse)
```

decision tree test mse: 4.515



Pruning does indeed help improve the MSE.

## Problem 2(d)

```
In [12]: from sklearn.ensemble import BaggingRegressor

bag_reg = BaggingRegressor(DecisionTreeRegressor(random_state=0), random_state=0)
bag_reg.fit(X_train, y_train);
```

```
In [13]: mse = mean_squared_error(y_test, bag_reg.predict(X_test))
print("bagged decision tree test mse:", mse)
```

bagged decision tree test mse: 2.6994652600000006

```
In [14]: feature_importances = np.mean([tree.feature_importances_ for tree in bag_reg.estimators_])
idxs = np.argsort(feature_importances)[::-1]

for col, imp in zip(X.columns[idxs], feature_importances[idxs]):
    print("%15s" %col, " |  %3f" %imp)

print("Top 3 features are:", X.columns[idxs][0], ", ", X.columns[idxs][1], ", ", X.columns[idxs][2])
```

Price		0.292875
ShelveLocMedium		0.215514
Age		0.121255
CompPrice		0.095484
Advertising		0.077952
Income		0.050261
ShelveLocBad		0.040722
Population		0.036205
ShelveLocGood		0.033244
Education		0.027782
Urban		0.006050
US		0.002655

Top 3 features are: Price ShelveLocMedium Age

## Problem 2(e)

```
In [15]: from sklearn.ensemble import RandomForestRegressor

rf_reg = RandomForestRegressor(random_state=0)
rf_reg.fit(X_train, y_train);
```

```
In [16]: mse = mean_squared_error(y_test, rf_reg.predict(X_test))
print("random forest test mse:", mse)
```

random forest test mse: 2.263103768200001

```
In [17]: feature_importances = np.mean([tree.feature_importances_ for tree in rf_reg.estimators_])
idxs = np.argsort(feature_importances)[::-1]

for col, imp in zip(X.columns[idxs], feature_importances[idxs]):
    print("%15s" %col, " |  %3f" %imp)

print("Top 3 features are:", X.columns[idxs][0], ", ", X.columns[idxs][1], ", ", X.columns[idxs][2])
```

Price	0.290875
ShelveLocMedium	0.195175
Age	0.105819
CompPrice	0.098693
Advertising	0.067502
ShelveLocBad	0.064920
Income	0.052249
ShelveLocGood	0.043733
Population	0.038941
Education	0.031178
US	0.005607
Urban	0.005308

Top 3 features are: Price ShelveLocMedium Age

```
In [18]: ms = np.arange(1, 12)

for m in ms:

    rf_reg = RandomForestRegressor(random_state=0, max_features=m)
    rf_reg.fit(X_train, y_train);

    mse = mean_squared_error(y_test, rf_reg.predict(X_test))

    print("m: %2d" % m, " | MSE: %.3f" % mse)
```

```
m: 1 | MSE: 3.131
m: 2 | MSE: 2.800
m: 3 | MSE: 2.484
m: 4 | MSE: 2.439
m: 5 | MSE: 2.350
m: 6 | MSE: 2.242
m: 7 | MSE: 2.185
m: 8 | MSE: 2.288
m: 9 | MSE: 2.316
m: 10 | MSE: 2.240
m: 11 | MSE: 2.328
```

As  $m$  is increased, we increase model flexibility up to a certain point, after which the model begins to have trouble identifying the signal in the data, as seen by the increase in MSE after about  $m \sim 7$ .

## Problem 3 (Chapter 8, Exercise 10)

```
In [19]: hitters = pd.read_csv("data/Hitters.csv")

print(hitters.info())
print(hitters.head())

# drop hitter names
hitters = hitters.drop("Unnamed: 0", axis=1)

# fix categorical columns
for col in ["League", "Division", "NewLeague"]:
    hitters[col] = hitters[col].astype('category').cat.codes
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 322 entries, 0 to 321
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Unnamed: 0            322 non-null   object
1   AtBat                 322 non-null   int64
2   Hits                 322 non-null   int64
3   HmRun                 322 non-null   int64
4   Runs                 322 non-null   int64
5   RBI                  322 non-null   int64
6   Walks                 322 non-null   int64
7   Years                 322 non-null   int64
8   CAtBat                322 non-null   int64
9   CHits                 322 non-null   int64
10  CHmRun                 322 non-null   int64
11  CRuns                 322 non-null   int64
12  CRBI                  322 non-null   int64
13  CWalks                322 non-null   int64
14  League                 322 non-null   object
15  Division               322 non-null   object
16  PutOuts                322 non-null   int64
17  Assists                322 non-null   int64
18  Errors                 322 non-null   int64
19  Salary                 263 non-null   float64
20  NewLeague              322 non-null   object
dtypes: float64(1), int64(16), object(4)
memory usage: 53.0+ KB
None

```

	Unnamed: 0	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	\
0	-Andy Allanson	293	66	1	30	29	14	1	293	
1	-Alan Ashby	315	81	7	24	38	39	14	3449	
2	-Alvin Davis	479	130	18	66	72	76	3	1624	
3	-Andre Dawson	496	141	20	65	78	37	11	5628	
4	-Andres Galarrraga	321	87	10	39	42	30	2	396	

	CHits	...	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists	Errors
0	66	...	30	29	14	A	E	446	33	20
1	835	...	321	414	375	N	W	632	43	10
2	457	...	224	266	263	A	W	880	82	14
3	1575	...	828	838	354	N	E	200	11	3
4	101	...	48	46	33	N	E	805	40	4

	Salary	NewLeague
0	NaN	A
1	475.0	N
2	480.0	A
3	500.0	N
4	91.5	N

[5 rows x 21 columns]

## Problem 3(a)

```

In [20]: hitters = hitters.dropna()
         hitters.info()

```



```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 263 entries, 1 to 321
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   AtBat           263 non-null   int64
1   Hits            263 non-null   int64
2   HmRun           263 non-null   int64
3   Runs            263 non-null   int64
4   RBI             263 non-null   int64
5   Walks           263 non-null   int64
6   Years           263 non-null   int64
7   CAtBat          263 non-null   int64
8   CHits           263 non-null   int64
9   CHmRun          263 non-null   int64
10  CRuns           263 non-null   int64
11  CRBI            263 non-null   int64
12  CWalks          263 non-null   int64
13  League          263 non-null   int8
14  Division        263 non-null   int8
15  PutOuts         263 non-null   int64
16  Assists         263 non-null   int64
17  Errors          263 non-null   int64
18  Salary          263 non-null   float64
19  NewLeague       263 non-null   int8
dtypes: float64(1), int64(16), int8(3)
memory usage: 37.8 KB

```

```
In [21]: hitters.head()
```

```
Out[21]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks
1	315	81	7	24	38	39	14	3449	835	69	321	414	...
2	479	130	18	66	72	76	3	1624	457	63	224	266	...
3	496	141	20	65	78	37	11	5628	1575	225	828	838	...
4	321	87	10	39	42	30	2	396	101	12	48	46	...
5	594	169	4	74	51	35	11	4408	1133	19	501	336	...

```
In [22]: hitters["Salary"] = np.log(hitters["Salary"])
hitters.head()
```

```
Out[22]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks
1	315	81	7	24	38	39	14	3449	835	69	321	414	...
2	479	130	18	66	72	76	3	1624	457	63	224	266	...
3	496	141	20	65	78	37	11	5628	1575	225	828	838	...
4	321	87	10	39	42	30	2	396	101	12	48	46	...
5	594	169	4	74	51	35	11	4408	1133	19	501	336	...

## Problem 3(b)

```
In [23]: X = hitters.drop("Salary", axis=1)
y = hitters["Salary"]

X_train = X.iloc[:200, :]
y_train = y.iloc[:200]
X_test = X.iloc[200:, :]
y_test = y.iloc[200:]
```

## Problem 3(c) & 3(d)

```
In [24]: from sklearn.ensemble import GradientBoostingRegressor

lrs = np.logspace(-4, 0, 17)
train_mse_hist = []
test_mse_hist = []


for lr in lrs:

    gb_reg = GradientBoostingRegressor(random_state=0, learning_rate=lr, n_estimators=100)
    gb_reg.fit(X_train, y_train)

    test_mse = mean_squared_error(y_test, gb_reg.predict(X_test))

    train_mse_hist.append(gb_reg.train_score_[-1])
    test_mse_hist.append(test_mse)
```

```
In [25]: plt.plot(lrs, train_mse_hist, marker="o", label="Training Loss")
plt.plot(lrs, test_mse_hist, marker="o", label="Testing Loss")
plt.xlabel("Shrinkage Parameter")
plt.ylabel("MSE")
plt.xscale("log")
plt.legend()
plt.grid()
```



## Problem 3(e)

```
In [26]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

mse = mean_squared_error(y_test, lin_reg.predict(X_test))

print("gradient boosting test MSE: %.3f" % np.min(test_mse_hist))
print("linear regression test MSE: %.3f" % mse)

gradient boosting test MSE: 0.207
linear regression test MSE: 0.492
```

## Problem 3(f)

```
In [27]: gb_reg = GradientBoostingRegressor(random_state=0, learning_rate=lrs[np.argmax(
                                                    n_estimators=1000, criterion="mse")
gb_reg.fit(X_train, y_train)

idxs = np.argsort(gb_reg.feature_importances_)[::-1]

for col, imp in zip(X.columns[idxs], gb_reg.feature_importances_[idxs]):
    print("%15s" %col, " |  %3f" %imp)

print("Top 3 features are:", X.columns[idxs][0], ", ", X.columns[idxs][1], ", ",
```

CAtBat		0.558630
CHits		0.095260
CRuns		0.053545
AtBat		0.049422
Walks		0.045620
CRBI		0.038976
CWalks		0.034677
CHmRun		0.029765
Years		0.023699
Hits		0.022652
RBI		0.016421
PutOuts		0.011897
Runs		0.010749
Errors		0.003875
Assists		0.002180
HmRun		0.001302
NewLeague		0.000923
League		0.000211
Division		0.000196

Top 3 features are: CAtBat CHits CRuns

## Problem 3(g)

```
In [28]: bag_reg = BaggingRegressor(DecisionTreeRegressor(random_state=0), random_state=0)
bag_reg.fit(X_train, y_train);

mse = mean_squared_error(y_test, bag_reg.predict(X_test))

print("gradient boosting test MSE: %.3f" % np.min(test_mse_hist))
print("bagging test MSE: %.3f" % mse)
```

gradient boosting test MSE: 0.207  
 bagging test MSE: 0.258

## Problem 4 (Chapter 8, Exercise 11)

```
In [29]: caravan = pd.read_csv("data/Caravan.csv")
```

### Problem 4(a)

```
In [30]: print(caravan.head())

# fix categorical column
enc = OneHotEncoder(categories=[["Yes", "No"]], sparse=False)
caravan["Purchase"] = enc.fit_transform(caravan["Purchase"].to_numpy().reshape(
```

	MOSTYPE	MAANTHUI	MGEMOMV	MGEMLEEF	MOSHOOFD	MGODRK	MGODPR	MGODOV	\
0	33	1	3	2	8	0	5	1	
1	37	1	2	2	8	1	4	1	
2	37	1	2	2	8	0	4	2	
3	9	1	3	3	3	2	3	2	
4	40	1	4	2	10	1	4	1	

	MGODGE	MRELGE	...	APERSONG	AGEZONG	AWAOREG	ABRAND	AZEILPL	APLEZIER	\
0	3	7	...	0	0	0	1	0	0	
1	4	6	...	0	0	0	1	0	0	
2	4	3	...	0	0	0	1	0	0	
3	4	5	...	0	0	0	1	0	0	
4	4	7	...	0	0	0	1	0	0	

	AFIETS	AINBOED	ABYSTAND	Purchase
0	0	0	0	No
1	0	0	0	No
2	0	0	0	No
3	0	0	0	No
4	0	0	0	No

[5 rows x 86 columns]

```
In [31]: X = caravan.drop("Purchase", axis=1)
y = caravan["Purchase"]

X_train = X.iloc[:1000, :]
y_train = y.iloc[:1000]
X_test = X.iloc[1000:, :]
y_test = y.iloc[1000:]
```

### Problem 4(b)

```
In [32]: from sklearn.ensemble import GradientBoostingClassifier

gb_clf = GradientBoostingClassifier(random_state=0, learning_rate=0.01, n_estimators=100)
gb_clf.fit(X_train, y_train)

idxs = np.argsort(gb_clf.feature_importances_)[::-1]

print("First ten features by importance:")
for col, imp in zip(X.columns[idxs][:10], gb_clf.feature_importances_[idxs][:10]):
    print("%15s" %col, " | " %3f" %imp)

print("Top 3 features are:", X.columns[idxs][0], ", ", X.columns[idxs][1], ", ", X.columns[idxs][2])
```

First ten features by importance:

PPERSAUT		0.074507
MOSTYPE		0.065550
ABRAND		0.056751
MGODGE		0.052928
MKOOPKLA		0.047574
MOPLHOOG		0.045789
MBERMIDD		0.040154
MGODPR		0.032920
PPLEZIER		0.031407
PBRAND		0.031169

Top 3 features are: PPERSONAUT MOSTYPE ABRAND

## Problem 4(c)

```
In [33]: y_pred = gb_clf.predict_proba(X_test)[:, 1]
y_pred = y_pred > 0.2
```

```
In [34]: from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[4336  197]
 [ 251   38]]
```

```
In [35]: tn, fp, fn, tp = cm.ravel()

gb_precision = tp / (tp + fp)

print("gradient boosting precision: %.3f" % gb_precision)
```

gradient boosting precision: 0.162

```
In [36]: from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression(random_state=0, max_iter=1E6)
log_reg.fit(X_train, y_train)

y_pred = log_reg.predict_proba(X_test)[:, 1]
y_pred = y_pred > 0.2

cm = confusion_matrix(y_test, y_pred)
print(cm)

tn, fp, fn, tp = cm.ravel()

log_reg_precision = tp / (tp + fp)

[[4293  240]
 [ 239   50]]
```

```
In [37]: print("gradient boosting precision:  %.3f" % gb_precision)
print("logistic regression precision: %.3f" % log_reg_precision)

gradient boosting precision:  0.162
logistic regression precision: 0.172
```

Gradient boosting does slightly better on the specified precision metric.